

Infrastructure as Code for Cybersecurity Training

Rui Pinto
University of Porto
Porto, Portugal

up201806441@up.pt

<https://orcid.org/0009-0003-5362-7819>

Rolando Martins
University of Porto
Porto, Portugal

rmartins@fc.up.pt

<https://orcid.org/0000-0002-1838-1417>

Carlos Novo
University of Porto
Porto, Portugal

carlos.novo@fe.up.pt

<https://orcid.org/0009-0003-0094-5565>

Abstract—An organization’s infrastructure rests upon the premise that cybersecurity professionals have specific knowledge in administrating and protecting it against outside threats. Without this expertise, sensitive information could be leaked to malicious actors and cause damage to critical systems.

In order to facilitate this process, the presented work addresses the use of Infrastructure as Code (IaC) and DevOps to automate the deployment of cyber ranges. An approach closely related to virtualization and containerization as the code’s underlying infrastructure helps lay down this burden. Notably, placing emphasis on using IaC tools like Ansible eases the process of configuration management and provisioning of a network. Lastly, several up-to-date vulnerabilities that are constantly messing with the lives of individuals and organizations are explored, most related to Privilege Escalation, Remote Code Execution attacks, and Incident Forensics, allowing the improvement of skills concerning Red team and Blue team scenarios. In short, one of the key takeaways of this work is contributing to better prepare specialists in ensuring that the principles of the National Institute of Standards and Technology (NIST) Cybersecurity Framework hold, namely: prevent, detect, mitigate, and recover.

Index Terms—Infrastructure as Code, DevOps, Cyber Range, Cybersecurity, Virtualization

I. INTRODUCTION

Preparing cybersecurity professionals to better respond to incidents using cyber ranges is costly due to the infrastructure complexity these setups may require and because the development of new scenarios is mostly a manual process. With paper and pencil training, it is possible to go over a vulnerable scenario, but details on how systems respond to incidents are often lost to the trainee. Moreover, many cyber range deployments are based on old case-driven methodologies that rely on hardware [7] and preconfigured virtualization through Virtual Machines (VMs) [13], [4], [8]. Containerization is starting to emerge [12], [11] as a more lightweight approach, but configuration management and deployment of these containers is often very specific to each implementation, turning the solution unscalable. For this reason, expanding the current cyber range platforms to diversify scenarios continues to be an issue.

Another possible development regarding cyber ranges is building networks that include containers and VMs [5]. This allows exploration of both generic and kernel vulnerabilities as, contrary to what happens with containers, VMs do not share the kernel with the host system.

II. MOTIVATION

With all the above situations in mind, the presented work focuses on developing and deploying a cyber range framework and exploring the creation of complex scenarios that the cybersecurity workforce will find helpful in refining their skills. There is a clear need to evolve this type of research so that the development costs are significantly reduced by taking advantage of virtualization techniques. Lastly, there is a tremendous need to familiarize ourselves with current attack scenarios, such as Log4j and Ransomware, and even old events like the Shellshock vulnerability so that the mistakes that happened in the past do not get repeated in the future.

III. GOAL

This work addresses the deployment automation of the software used for cybersecurity training using an approach based on Infrastructure as Code and DevOps for networking and the relying infrastructure used by these scenarios. Several vulnerabilities will be explored related to Remote Code Execution, web applications, Privilege Escalation, and forensics, which are associated with the daily threats companies face. The ultimate goal is to build a set of playbooks that will automatically deploy, configure and provision container-based environments in a reasonable amount of time and use fewer resources in terms of funds and infrastructure so that the deployment process happens more efficiently.

Besides, running an entire enterprise-level network in a single computer or in the cloud is possible due to the scenarios’ containerized nature. Notions regarding the safety of the training system are taken into account to ensure that no actual harm reaches the host computer but the sandboxed environment.

IV. RELATED WORK

A. Hardware-based Cyber Ranges

According to Ferguson *et al.* [7], the National Cyber Range (NCR), closely tied to the American Department of Defense, provides a “*unique environment for cybersecurity testing throughout the program development life cycle using unique methods to assess resiliency to advanced cyberspace security threats.*” The NCR works as an Internet-like environment supported by a multitude of VMs and physical hardware. Scenarios are deployed to perform tests that should not occur on open operational networks due to potentially catastrophic consequences caused by the execution of malicious payloads.

It features traffic generation techniques, several types of vulnerability scanning, exploitation, and data capturing tools. As expected, the main focus of this project is purely military, meaning there are few details on the internals of the cyber range.

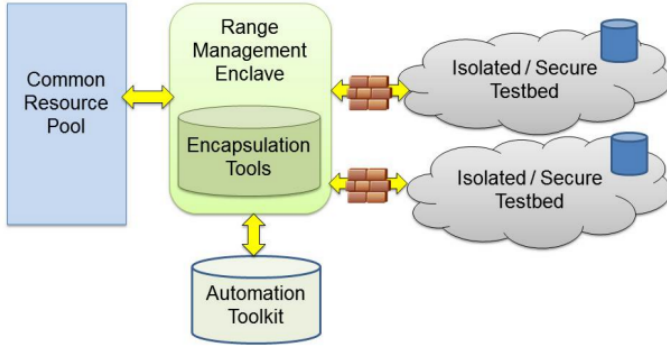


Fig. 1: NCR Core Capabilities [7].

As depicted in Fig. 1, a firewall is placed between the “Isolated Testbeds” and the “Range Management Enclave.” The latter consists of encapsulation tools and an automation tool kit that provisions resources from a “Common Resource Pool.” Physical Layer 1 switching ensures isolation concerning the low-level communication protocol stack.

The sequence of actions required to execute tests starts with assigning hardware and software resources from the “Common Resource Pool.” Afterward, the provisioning process includes using the Layer 1 switch to isolate the selected resources from all the other NCR assets. At this point, the systems under test are installed, and the final network state is matched against the initial expectations. Then, tests are performed, and results are collected. Lastly, hardware is sanitized, ensuring no remnants related to the test, and made available back in the “Common Resource Pool.”

Gustafsson *et al.* [9] proposes CRATE, a cyber range heavily relying on a dedicated hardware platform and virtualization. The high-level architecture of this system is featured with a set of *virtualization servers* that house the emulated environments, a *control plane* used for management tasks, and the *event plane* for systems where training sessions are executed. Inside each virtualization server, a customized Linux-based operating system called CrateOS was placed. Among many other features, it contains a system service named NodeAgent that handles communication between the Core Application Programming Interface (API), which connects to the Application Layer, the Database layer, and the VMs. It automates the deployment and configuration activities of the cyber range. The network in the event plane uses Software Defined Networking (SDN) to facilitate automated configuration and emulation of the networks. Besides, virtual network segments, VXLANs, are used to support many emulated networks. Furthermore, CRATE *Exercise Control (CEC)* is a tool used to set up and manage training sessions. *SVED (Scanning, Vulnerabilities, Exploits, and Detection)* is used to automate experiments

and training scenarios. It consists of several modules: one linked to vulnerability data and automatic scans performed with OpenVAS, and others related to designing attack graphs, executing them, and generating reports. CRATE allows connecting any hardware device in the emulated environments to conduct experiments with hardware-based security solutions. It is featured with traffic generation tools and data collection tools, using *tcpdump* and *Snort*.

B. VM-based Cyber Ranges

Most current *state-of-the-art* focuses on cyber ranges based on VMs. Pham *et al.* [13] proposes a system, CyRIS (Cyber Range Instantiation), that automatically prepares and manages cyber ranges for cybersecurity training based on custom specifications. CyRIS is part of CyTrONE [4], a training framework that facilitates training activities, providing an open-source set of tools that automate the training content generation. It also integrates with a Learning Management System, Moodle.

CyRIS is a module that takes a configuration input file following the Yet Another Markup Language (YAML) format and a base image under the format used for Kernel-based Virtual Machine (KVM) virtualization, creating the desired environment according to the provided description. This base image contains a set of pre-installed operating systems and several basic system configurations (hostname, SSH keys, IP addresses). Later on, a master node running the CyRIS service processes the description file and allocates VMs that will be assigned to other hosts of the same LAN network, as observed in Fig. 2.




Fig. 2: CyRIS’s Architecture [13].

CyRIS comprises five key features that play important roles in its architecture:

- *System Configuration*, which not only involves basic system configuration but also managing user accounts and modifying firewall rules.

- *Tool Installation*, which is essential for the penetration testing work.
- *Incident Emulation*, as per the ability to launch actual incidents. It consists of attack emulation, traffic capture, and malware emulation.
- *Content Management* consisting of copying content into the cyber range, executing scripts, and generating logs.
- *Clone Management*, which considers the defined network topology for the VMs and the inherent isolation between them.

Beuran *et al.* [4] presents CyTrONE that follows the architecture presented in Fig. 3, where CyRIS is also represented.




Fig. 3: CyTrONE's Architecture [4].

The *Training Management* module is based on user inputs and the training database, which includes training scenarios (VM base images), security incidents, and vulnerability information. This module is responsible for creating the input files related to *content description* and a *cyber range description* defining the training's content and activity. The *Content Description Processing* module converts the content description to a format named SCORM, which is widely used in the e-learning industry and, therefore, understandable by Moodle. The adoption of Moodle is related to educational purposes and follows a Q&A approach, as questions related to the posed challenge will be presented on this platform.

Jiang *et al.* [10] mentions a particularly interesting VM-based type of cyber range, Pandora, which is intentionally incompatible with enterprise systems to reduce the risk of attack propagation into the infrastructure. It proposes a system suitable for automated testing of exploits and result collection, keeping security concerns related to the sandboxed environment in mind by considering vulnerabilities on VMware Fusion (CVE-2015-2337) and Venom (CVE-2015-3456) that allowed VM escape, thereby causing damage to host systems.

Pandora's architecture runs under a VM with an operating system that introduces some incompatibility with the "Generic Operating System" to, as mentioned before, introduce an intentional inconsistency with regards to damage propagation outside the testing environment. The "Vulnerable Binary Manager" is used to execute vulnerable binaries within the secure environment, use exploits against the vulnerable binary and record the effect of such exploitations. A "Vulnerable Binary" is a file containing a set of defined vulnerabilities that automated tools will exploit. Notice that this binary should only be able to be executed within the "Secure OS." The "Vulnerability Manager" is an API that handles communication with the cyber range by sending exploits and receiving responses from the "Vulnerable Binary Manager." The "Automated Cyberse-

curity Tool(s)" generates exploits against vulnerable binaries and is not present in the secure operating system to assure simplification. Examples include fuzzing tools, such as Fuzzer and American Fuzzy Lop, to generate crash strings for simple binary files vulnerable to buffer overflows. Later on, rex, an automated exploit engine, exploits the target binary using the above-mentioned crash string obtained from the fuzzing tool, generating a Proof of Vulnerability (POV) that is later on sent to the "Vulnerability Manager" and received by the "Vulnerability Binary Manager" inside the cyber range VM.

C. Container-based Cyber Ranges

Container-based cyber ranges are frequently linked to reduced resource consumption compared to VM-based scenarios, mainly because the container's resources are shared with the host, causing a lower overhead. This same overhead is even lower compared to scenarios full of virtual instances since the CPU and memory usage compared to VM-based scenarios is much lower. This section intends to elaborate on container-based cyber ranges and their specific details.

Perrone *et al.* [12] brings forward the *Docker Security Playground*, a microservices-based approach to building complex network infrastructures tailored to study network security. These microservices are based on Docker. Likewise, it offers an API enabling further development on the lab scenarios. This project uses Docker-compose to manage scenarios' start and stop procedures. It uses a *Docker Image Wrapper* which defines a standard notation for Docker labels to provide custom configurations for the base Docker image of each scenario. The Docker images used within the project are placed at DockerRegistry Hub.

Related to the *Docker Security Playground*, Caturano *et al.* [5] designed another container-based cyber range built upon the *Docker Security Playground*. This project tackles the fact that several vulnerabilities related to the kernel cannot be explored in Docker-based environments because containers share the Linux kernel with the Docker host. To overcome this problem, Caturano *et al.* explores scenarios based on containers and VMs, providing a hybrid environment where cybersecurity exercises are deployed. For this, it makes use of "macvlan" interface drivers to create a bridge between Docker containers and VMs. Similarly, Acheampong *et al.* [1] mixes the concept of cloud deployments based on VMs with Docker containers hosting packaged applications.

Closely related to the education side, Thompson *et al.* [15] refers to another training framework, Labtainers, which relies on Docker containers featured with an automatic assessment of students, as lab data is collected and automatically sent to an instructor upon completion via email or a Learning Management System, creating the possibility of analyzing the experiential learning efficacy of the exercise. Specific properties of the laboratory activities are randomized so that each student works on a different scenario in terms of configurations. This randomization is achieved by defining symbols within the source code and data files part of the lab, which are replaced with student-specific values upon lab startup, for instance,

the buffer size related to a buffer overflow vulnerability. This concept is further explored in the next section.

D. Randomization

Developing a scenario for a cyber range is mainly a manual process. An example that corroborates this statement is SEED Labs [6], where many of Labtainers' laboratories are based [15], where several scenarios based on software security, web security, system security, mobile security, network security, cryptography, and blockchain are made available to trainees using both VMs and containers. More than 80 universities use SEED Labs, which expresses close bounds to education.

Consequently, there is a need to address some randomization of the developed scenarios as they essentially turn out to be static once created. As so, this section elaborates more on the cyber ranges, both VM-based and container-based, that take into account randomization.

Starting with Schreuders *et al.* [14] that proposes a VM-based cyber range, SecGen, developed in Ruby that introduces randomization. It is suited for both educational lab usage and Capture The Flag (CTF) challenges. One of the main concerns is addressing the sluggish pace associated with the manual configuration of hacking scenarios, which is not practical at scale. This cyber range focuses on a CTF-style type of challenge, where solving the proposed challenge results in discovering a secret flag. The introduced randomization is characterized as follows:

- *Selection*: randomized selection of the operating system, network configurations, services, system configurations, and vulnerabilities to be used.
- *Parameterisation*: that entails system elements should be configurable, for instance, the strength of a user account password.
- *Nesting*: data generation should be combined/nested randomly.

The description of the system greatly depends on the eXtensible Markup Language (XML) specification language, which states the details related to the configuration of the network, available vulnerabilities, services, users, and content and applies logic for randomizing the scenario. It uses Puppet and Vagrant to provision the VMs. A critical takeaway idea of this project is its highly modular structure and the use of vulnerabilities and associated exploits provided by the Metasploit Framework. The SecGen running process is composed of two stages:

- *First Stage*: is where all the scenario Ruby modules are read, randomization steps are applied, and the Puppet modules are deployed. At last, a Vagrant file is created, which describes the entire scene, according to the steps mentioned.
- *Second Stage*: leverages Vagrant to generate and provision the VMs.

Currently, SecGen counts over 100 modules: data generation modules, encoder modules, providing various encryption and encoding methods, service modules, providing a wide range of

secure services, utility modules, allowing various system configurations, and vulnerability modules, concerning vulnerable services.

Consequently, Nakata *et al.* [11] proposes a Directed Acyclic Graph (DAG) based cyber range, CyExec, with randomization techniques in mind, using Docker containers. This article claims CyExec outperforms the SecGen VM-based scenario [14] generator, consuming 1/3 memory, having 1/4 CPU load, and 1/10 of storage usage, primarily since it uses containers instead of VMs.

The concept of randomization here takes the form of a graph, being each milestone a vertex and each scenario an edge, meaning an attack consists of different types of vulnerabilities that achieve the same outcome. As a result, the trainee experiences several distinct situations. Since the attack is directed towards the final target and there is no going back to a previous milestone or looping back to the same vertex, this graph is considered a DAG.

Fig. 4 shows the structure followed by CyExec. With several Docker-compose files, randomization is assured because it allows switching between which *Dockerfiles* are used when setting up a scenario. A *Dockerfile* works as just another "edge" to reach a "vertex," leading us to the fact that different vulnerabilities are introduced into the system according to the selected *Dockerfile*.




Fig. 4: CyExec structure [11].

For example, consider a scenario where Metasploitable2, a purposely vulnerable system, is available in the CyExec testbed. Several vulnerable applications can be found in this machine, such as vsftpd, PHP, Samba, and PostgreSQL, among others. Scenarios can be swapped and selected randomly if different vulnerabilities or attack techniques are considered.

E. Summary

As mentioned, current cyber range problems are related to high-cost infrastructure relying too much on VMs, and details of these kinds of open-source scenarios rarely reference up-to-date vulnerabilities and attacks. Container-based solutions are starting to appear, but they heavily rely on custom cyber range description files, using the YAML, XML, or, sometimes, even the JavaScript Object Notation (JSON) format. There are cases where IaC tools are not being used, hampering scalability concerns. In such cases, a custom tool was typically designed to parse and take action on the aforementioned customized scenario description files. Randomization is clearly lacking in some cyber range scenarios, which is understandable given that developing these testing environments is a slow manual process.

TABLE I: Comparison of Cyber Ranges.

	IaC	Randomization	Local & Cloud Deployments	Containerization	Enterprise-level Scenarios	Linux & Windows Scenarios	Open-source
<i>CyExec</i>	◐	●	●	●	◐	◐	◐
<i>Pandora</i>	◐	●	◐	◐	◐	◐	◐
<i>CyRIS/CyTrONE</i>	◐	●	●	◐	◐	◐	●
<i>NCR</i>	-	●	◐	◐	◐	-	◐
<i>SmallWorld</i>	◐	◐	●	◐	●	●	◐
<i>Leaf</i>	●	●	●	◐	●	◐	◐
<i>CRACK</i>	◐	●	◐	◐	●	◐	●
<i>SEED Labs</i>	◐	◐	◐	◐	◐	◐	●
<i>Labtainers</i>	◐	●	●	◐	●	◐	●
<i>CRATE</i>	-	●	◐	◐	●	◐	◐
<i>DSP</i>	◐	◐	◐	●	●	◐	●
<i>SecGen</i>	●	●	◐	◐	●	◐	●
<i>Proposed Solution</i>	●	●	●	●	●	●	●

Table I presents a high-level overview of the main features supported by some of the most relevant cyber ranges, finishing with the proposed solution's aim. Several marks were given according to each cyber range. As a way to demystify some classifications, for the *IaC* column, half a circle considers approaches that used customized descriptions to deploy scenarios or solutions only relying on Docker or Docker-compose without a standard tool. For the *Randomization* column, tools that provided some randomization of traffic generation, network, system, and accounts' configurations or vulnerabilities present in the scenario were considered a full circle. Regarding the *Local & Cloud Deployments* and *Linux & Windows Scenarios*, half a circle was considered for cases where only one of the features was present. About the *Containerization* column, half a circle was considered for scenarios that combined containers and VMs in separate scenarios. Concerning the *Enterprise-level Scenarios* column, full circles were considered for networks with a wide variety of services ranging from firewalls, internal networks, Intrusion Detection Systems (IDS), and mail servers, among others. At the same time, half-circles were intermediate representations of enterprise-level networks. Finally, the *Open-source* column considers scenarios available to the general public.

This framework addresses every column of Table I, which is not achieved by any other framework. Even in cyber range frameworks such as *CyExec*, which is complete in terms of the mentioned features, a possible idea would be to extend the development. Unfortunately, not all frameworks are open-source, and some lack community support. Instead, a new framework with custom functionalities was created using the technology stack of the development team's choice.

V. DEVELOPED WORK

A. Architecture

The scenario construction process using Docker containers targeted enterprise-level networks. As such, corporate environments normally subdivide networks into three different main sections:

- **External Network** refers to the public internet where machines are not controlled by the organization. As such, risk modeling activities should be taken into account in

order to evaluate the risk and the probability specific threats and attack scenarios pose to the internals of the organization. With this, according to the organization's budget, decisions on which security measures to place in the company's network are considered and may include systems like IDS, Intrusion Prevention Systems (IPS), Firewalls, Antivirus, among others.

- **Internal Network** which contains the protected machines of an organization, such as internal databases and services only available to the company's employees and not to the general public.
- **Demilitarized Zone (DMZ)**, which is a network that protects the company's internal network and is targeted with untrustworthy traffic. It includes services available to the public and sits between the *External Network* and the *Internal Network*. It generally includes web servers, Domain Name System (DNS) servers, among others.

This project focuses on these three distinct types of networks and considers several network services that people would typically see on enterprise networks.

The network architecture presented in Fig. 5 shows the services available on every Linux scenario, except for Windows-based scenarios, which slightly differ from this schema. As shown, Ansible appears as the tool responsible for configuring and provisioning the entire network.

B. Ansible Architecture

Three different playbooks include all the developed scenarios. The first is explicitly used in Linux-based scenarios, representing most designed challenges. The second is used for the Windows Ransomware scenario, and the last for the Windows Active Directory (AD) scenario. On each playbook, the first step is to delete stale Docker containers from previous running scenario executions, as shown in Listing 1.

```

- hosts: localhost
  pre_tasks:
    - name: Remove Stale Containers
      ansible.builtin.include_tasks:
        teardown.yml

```

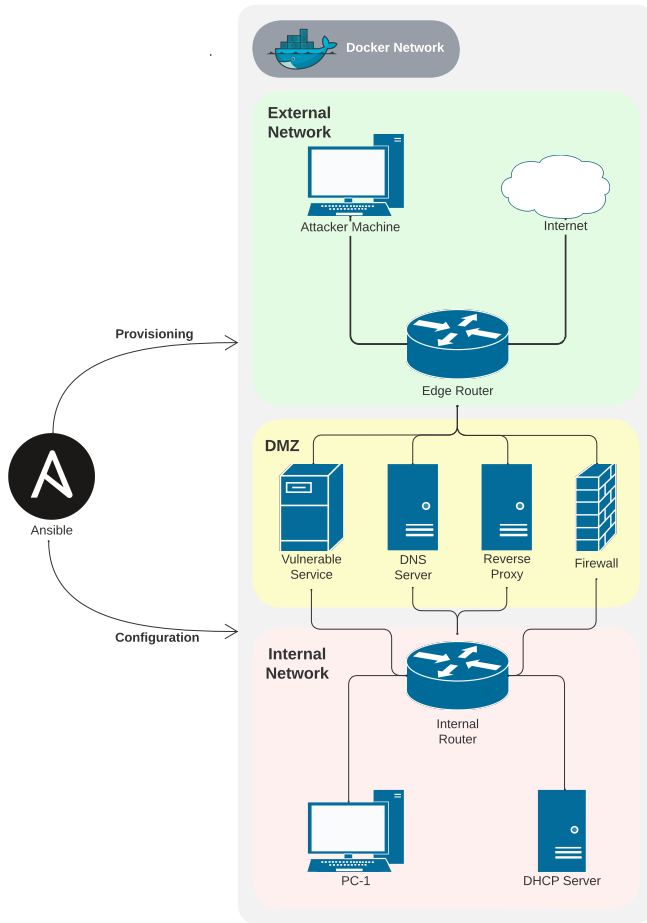


Fig. 5: Template Network Architecture.

```
loop: "{{ machines + vulnerables.
machines }}"
loop_control:
loop_var: pc_info
```

Listing 1: Removal of Stale Containers.

Essentially, for every machine object passed, the contents of the *teardown.yml* file are run. This uses the *community.docker.docker_container* module that is built-in in Ansible and removes the container under a given name.

C. Ansible Groups and Inventory

Every machine belongs to a group, by default in Ansible, the *all* group. Nonetheless, other groups and respective members were defined in the so-called Ansible Inventory, as presented in Listing 2.

```
[routers]
[firewalls]
[external]
[internal]
→ [pcs]
→ [dhcp_servers]
```

```
[dmz]
→ [dns_servers]
→ [custom_machines] # Scenario's vulnerable
machines.
→ [reverse_proxies]
```

Listing 2: High-level View of Ansible Inventory.

Each word represents a group of one or more machines. Each group may have several child groups defined by their name, as it happens above, or by machines, represented by their Fully Qualified Domain Name or IP address. Groups themselves are very useful when restricting specific tasks per group. Then, some groups contain child groups, as happens with the *internal* and *dmz* groups. In the case of Windows-based scenarios, another group called *machine* is used and refers to the Docker container containing the Windows Vagrant box. Listing 2 is a very high-level view of how groups are organized within the project. A custom Python inventory script was created to allow the specification of custom variables across each group.

D. Generic Scenario Variables

For each playbook, a set of variables is always defined and corresponds to the generic structure of the network, as presented in Fig. 5. We start with the Docker images used across the workflow, their path, and the default image name in case none is specified.

```
general:
images:
- name: kali_test_img
path: ./attacker
- name: base_image
path: .
default_container_image_name: base_image
```

Listing 3: Ansible Variables - Docker Images.

As shown in Listing 3, two Docker images are used: *base_image* and *kali_test_img*. The former is an image derived from *node:lts-alpine* with some extra packages installed. The Alpine distribution was chosen due to its smaller size compared to other images. As a result, the *base_image* size is around 230MB. The *kali_test_img* is an image derived from the official *kalilinux/kali-rolling* Docker image. This image was extended to include the Xfce desktop environment, characterized by its low resource consumption and user-friendliness, as well as the *Virtual Network Computing* (VNC) package, which allows screen sharing and remote control from another device, meaning the computer screen, keyboard, and mouse are mapped from an external device to the device installed with VNC. Accessing port 6080 on the target machine makes it possible to obtain remote control over it, which will be later used in the scenarios. This Kali Linux image is especially suited for offensive tasks, and here the only concern was providing the trainee with a broad range of tools he could

use in a scenario. Therefore, the image's size is much larger (around 11GB) compared to the *base_image* used for common network services.

The second category of Ansible variables for machines belonging to the *all* group concerns Docker networks, according to the structure mentioned in Section V-A. The range of each network is defined, as well as the gateway address which points to the host machine. This is mandatory by Docker, as the host machine should always take part in each created virtual Docker network to forward packets from and to it. At last, the *random_byte* attribute points to a random byte that changes across each scenario execution and confers some degree of randomization, as for each new scenario execution, the networks' IP addresses will change.

Several attributes are specified for a machine's variables according to the logic of a Docker container. Starting with its name, the Docker image it uses, possible volumes (anonymous, named, or bind mounts), the groups the container belongs to, and published ports, meaning ports mapped between the Docker container and the host machine. Then, the container is assigned to one or more networks. Lastly, the DNS server location is specified. In this case, as the attacker machine is located in the external network, DNS queries are redirected to the edge router's network interface sitting in the external network so that these queries are later forwarded to the DNS server in the DMZ. This is achieved using *iptables* rules. For machines located inside the corporate network, DNS queries are sent directly to the DNS server sitting in the DMZ network without the need for any type of forwarding by the edge router. It is also important to mention other attributes that are also possible to be specified, namely the *devices* and *privilege* attributes, all having the same meaning as understood by Docker.

E. Custom Scenario Variables

After presenting how the standard setup for each scenario is organized, Listing 4 shows the structure of the scenario's custom variables, starting with an example of a DNS configuration.

```
dns:
- domain: example-domain.ui.com
  internal:
    machine: vuln_service
    network: dmz_net
  external:
    machine: edge_router
    network: external_net
```

Listing 4: Ansible Variables - DNS.

Here, a domain named *example-domain.ui.com* is presented along with *internal* and *external* specifications of it. This is related to two distinct DNS views that are defined. "Internal view" refers to devices in the internal or DMZ networks;

otherwise, they belong to the "external view." So, in the listing mentioned above, the *example-domain.ui.com* domain points to the *vuln_service* container located in the DMZ whenever devices in the "internal view" look for this domain. Devices in the "external view" point to the external network interface of the edge router. This means resolved DNS requests made by external machines will go through the edge router and are forwarded to the respective machine.

Furthermore, the set of variables concerning custom machines' Docker images has a similar format to the one presented in Listing 3. Still, the representation is a bit more flexible, allowing the specification of the name of the *Dockerfile* and arguments to be read in the Docker image creation process.

Then, in the vulnerable machines section, the situation is quite the same as presented for the generic machines. The only exception is the inclusion of an attribute that allows the specification of variables for each machine.

```
port_forwarding:
- destination_port: 443
  to_machine: reverse_proxy1
  to_network: dmz_net
  to_port: 443
```

Listing 5: Ansible Variables - Port Forwarding.

Then, Listing 5 references the port forwarding section especially relevant for external machines and how they can communicate with DMZ machines. The attributes meaning are as follows:

- *destination_port*: the incoming port on the edge router where packets will later be redirected.
- *to_machine*: the target machine to which packets reaching the *destination_port* will be forwarded to.
- *to_network*: the network where the target machine is placed.
- *to_port*: the destination port in the target machine where the edge router will redirect packets.

Some names may be misleading, such as *destination_port* and *to_port*. Still, they obey the convention used by *iptables*.

```
setup:
  machines:
- name: localhost
  setup: "{{ playbook_dir }}/scenarios/chessrs/setup/"
- name: attackermachine
  setup: "{{ playbook_dir }}/scenarios/chessrs/attacker_machine_setup/*.j2"
```

Listing 6: Ansible Variables - Setup Section.

Lastly, we have Listing 6, which provides information on where to find the setup instructions for the *localhost* and attacker's machines.

F. Ansible Roles & Network Services

The structure followed by Ansible uses a feature called "roles". A different role is used for every milestone in the network configuration. Ansible allows defining specific variables and tasks for each role, making grouping an entire workflow into separate roles straightforward to reuse in the development cycle. In the folder structure, a directory represents a single *role*. Inside it, specific tasks are defined. The following sections detail the tasks present for each role.

1) *Base Role*: The *base* role is responsible for the scenario's initial tasks:

- Start the Docker service.
- Building the scenario's Docker images, as presented in Listing 3.
- Create the Docker networks.
- Create generic scenario's Docker containers.
- Assign each created container to one or more Ansible groups.

2) *DHCP Role*: The *DHCP* role configures the DHCP servers. At first, the *dhcp* package is installed. Then a template configuration file is created using Jinja2 templates. At last, the *dhcp* service daemon is started.

Essentially, the DHCP lease is responsible for assigning an internal IP address with the last byte ranging from 64 to 127, pinpointing the router of the internal network as the gateway router, and updating the DNS server with the one placed in the DMZ network.

3) *Internal PCs Role*: The *internal PCs* role handles the behavior of machines inside the internal network. As such, it runs the following tasks:

- Install the DHCP client package.
- Ask for a DHCP lease to the DHCP server.
- Removes the automatically assigned IP address by Docker so that its only IP address is the one stated by the DHCP server.

4) *Internal Role*: The *internal* role is destined for the internal machines and the DHCP server. It simply configures each device's default route as the internal router's interface in the internal network. Every time a default gateway or static route is configured across the Ansible setup, the *iproute2* package, installed on each Docker image by default, is used. The *iproute2* package allows controlling and monitoring various aspects of networking in the Linux kernel, namely routing, tunnels, network interfaces, and traffic control, among others.

5) *DNS Role*: The *DNS* role is somewhat of a more complex role and is destined for DNS servers. It is responsible for running the following actions:

- Install the *bind* DNS server package.
- Copy the necessary template DNS configuration files to the DNS server container.
- Start the *named* DNS service.

Two Access Control Lists (ACLs) are created regarding the DNS configuration files. The *internal* deals with which machines stand in the internal and DMZ networks, and the *external* points to every other machine that is not part of the *internal* ACL, meaning external machines only. Afterward, a distinction on the IP addresses retrieved by resolved domains for internal and external machines is made, according to what was presented in Listing 4.

Essentially, an ACL named "exclude" is created and refers to the IP address of the edge router because it performs Network Address Translation (NAT) over specific packets coming from outside the organization's network. Then, the "internal" and "external" views are created, as mentioned above, which direct requests to different DNS zones according to the mapped domain. If the DNS query does not match any internal domain, the request will be forwarded to the 8.8.8.8 Google's public DNS server.

6) *Router Role*: The next role leads us to the router's configuration steps. Here, the steps vary according to the internal or edge router configuration steps.

The internal router takes a single action to configure its default gateway with the edge router's DMZ network interface, as this is the gateway that provides internet access to the network.

The edge router's task is to provide NAT to packets whose source matches the internal or DMZ networks. This is accomplished using the *MASQUERADE* jump of *iptables*. Lastly, a static route is added, specifying that packets destined for the internal network should be directed to the DMZ interface of the internal router.

7) *Custom Machines Role*: The *custom machines* role is quite similar to the *base* role except that it performs Docker image and container creation on the set of specific machines required by a scenario instead of the generic ones. Although both roles slightly differ, using just one role with some small conditionals would be possible. Still, the adopted approach allows more accessible future updates.

8) *DMZ Role*: The *DMZ* role is also quite simple. It targets DNS servers, custom machines, and reverse proxies sitting in the DMZ network. Two tasks are associated with this role: the static configuration route to the internal network, which directs packets to the internal router's interface on the DMZ network, and the configuration of the default gateway to access the internet, which points to the edge router's DMZ network interface.

9) *Reverse Proxies Role*: The *reverse proxies* role, as the name suggests, targets the reverse proxies present in the network. These services sit in front of the scenario's custom machines and forward client requests to those machines. Essentially, they allow establishing HTTPS connections with the client device, handling all the Secure Sockets Layer (SSL) certificate-related tasks from the Transport Layer Security (TLS) connection, and talking with the destination machine in the back of the reverse proxy using an HTTP connection. In simple words, it works as a middle agent.

The variables defined for the reverse proxy include a domain and information on the target machine. The current setup allows the configuration of several domains specified using Jinja2 templates. The reverse proxy continuously listens for connections at port 443, and according to the selected domain, it forwards the traffic to the appropriate target. If requests are made to port 80, they are redirected to port 443, meaning the HTTP connection gets upgraded to HTTPS.

Then, SSL certificates need to be taken into consideration. For this, a Certificate Authority (CA) was first created by defining an `openssl.cnf` file and the necessary folder structure, as well as generating the public and private keys associated with the CA.

The *reverse proxies* role generates a Certificate Signing Request (CSR), a specially formatted encrypted message sent from an SSL digital certificate applicant to a CA. The CA then takes the CSR and generates a public-key certificate signed by itself. Then, it removes the password from the private key associated with the newly issued public-key certificate and copies both files to the reverse proxy container. At last, it starts the NGINX service with the loaded configuration.

One crucial aspect of this configuration is the signing of public-key certificates by the CA. This entails that the created root CA has to be trusted by machines that will eventually access the domain linked to the digital certificate, in this case, the attacker machine. To achieve such setup, the CA's public-key certificate is loaded as trusted in the attacker machine both system-wide and in Firefox, as it will be later explained.

10) *Firewalls Role*: The *firewall* role focuses on the two existing routers which incorporate a firewall. During this explanation, the internal router's firewall will be referred to as the internal firewall, and the edge router's firewall as the external firewall.

Starting with the internal firewall, the role performs the following *iptables* actions:

- Set the forward chain's default policy to drop, meaning the internal router does not forward traffic by default.
- Already established connections or connections previously associated with existing ones to the internal network are accepted.
- New, previously established, or related connections from the internal network are also accepted.

Concerning the external firewall, this role performs the following *iptables* actions:

- **Generic Rules:**

- Set the forward chain's default policy to drop, meaning the external router does not forward traffic by default.
- Established connections or connections previously associated with existing ones to the internal or DMZ networks are accepted.
- Packets from the internal or DMZ networks are also accepted in the forward chain.

- **DNS Rules:**

- Set a "prerouting" chain rule in which TCP and UDP traffic reaching the edge router's port 53 will have its destination changed (DNAT) to the real DNS server sitting in the DMZ network and destination port 53.
- Forwarding traffic to the DNS server is accepted.
- A "postrouting" NAT rule is added to traffic whose destination is the DNS server.

- **Port Forwarding Rules:**

- Allow TCP and UDP forwarding according to the information provided in the example of Listing 5. The target machine and the target port number are considered in this regard.
- Accept TCP and UDP traffic in the "prerouting" chain according to the information provided in the example of Listing 5. The target machine and the destination port number are considered in this regard.
- A "postrouting" NAT rule for the traffic reaching the target machines, as in the example of Listing 5.

With both the internal and external firewalls, the allowed traffic is restricted from the devices externally placed with respect to the organization's network. Only certain services in the DMZ should be allowed external access, never devices from the internal network. On the other hand, connections from inside the corporate network are allowed. This configuration uses *iptables* to create a realistic firewall setup. Therefore, rules are not based on highly-complex logic like which domains an internal device tries to access and if they should be blocked, according to a blocklist of IP addresses and domains.

11) *Entry point Role*: The *entry point* role performs the necessary tasks to configure a particular machine, as defined at Listing 6. It creates the environment needed to run the setup scripts, which may include copying template files to the Docker container and then executing the entry point script. This role distinguishes when being conducted by the *localhost* machine or a different machine. For the *localhost*, the entry point script is run without any previous configuration. In the case of the other machines, the Jinja2 template setup files are first copied to the target container, and then the entry point script is run.

12) *Mesh Role*: The *mesh* role handles devices that need to join the Tailscale network, called *tailnet*. This network allows communication between each device that belongs to it. This is

useful when focusing on cloud deployments if, for instance, a connection is established on port 6080 in the attacker machine to be able to control it remotely and, as further detailed, in Windows-based scenarios to access the Windows Vagrant box using *Remote Desktop*.

This role's actions start by installing Tailscale and starting the *tailscaled* service. After this, the container is instructed to join a specific Tailscale network using an authentication key and by specifying a hostname for the machine. An authentication key allows the addition of new nodes to the Tailscale network without needing to sign in to the network. A reusable authentication key was created to connect multiple nodes to the network. Each time a new node joins the Tailscale network using this authentication key, it enters the group of Tailscale ephemeral nodes, which essentially refer to short-lived devices that are automatically removed from the network after a short period of inactivity and are immediately removed from the network in case they are instructed to do so. Also, the usage of the same hostname for a particular machine allows accessing it using a Tailscale feature called "MagicDNS" which essentially registers all the DNS names for the network's devices using the following schema: `[Device Hostname].[Network DNS Name]` The device's hostname was already mentioned above. By default, the network's DNS name is chosen by Tailscale upon the first usage but can be modified afterward. The "MagicDNS" configuration provides easy access to machines when they are not under the host's control. This will be useful in Section X.

VI. CUSTOM SCENARIOS

The set of custom scenarios involves three distinct cyber ranges:

- A Linux scenario that explores the Apache Log4j vulnerability (CVE-2021-44228).
- A Windows-based scenario that explores a Ransomware malware executable that encrypts a set of files.
- A Windows-based scenario that exposes a vulnerable Active Directory Domain Controller that provides a vast attack surface where the trainee can experiment with several attacks.

For each challenge, details on how to solve them will be revealed. Some of the intended solutions to get the secret flag are presented and, when available, unintended solutions.

A. Log4j Scenario

The Apache Log4j vulnerability started haunting the world during the last month of 2021. It was based on the Java-based logging package Apache Log4j and essentially allowed an attacker to execute code on a remote server, the so-called Remote Code Execution. The scope of machines this vulnerability targeted was enormous, and some put it on the same level as the most severe vulnerability along with *Heartbleed* and *Shellshock*. CVE-2021-44228 details which Log4j versions were affected and gives a brief insight into what the vulnerability is about. In simple words, an attacker that can control log messages may run arbitrary code by

means of a process called message lookup substitution.

1) *Scenario Construction*: This scenario is based on the Tier 2 Unified Hack The Box challenge and in the *SprocketSecurity* blog post [2], essentially reproducing a vulnerable version of the Ubiquiti UniFi network application dashboard. This works as an interface manager for all the hardware devices belonging to Ubiquiti's mesh network allowing the changing of several network-related configurations. To replicate the scenario, the Goofball222's GitHub UniFi Docker container repository version 6.4.54 was used with some tweaks in the *Dockerfile* suited for Alpine-based distributions by installing the `python3` package and removing the `JVM_EXTRA_OPTS=-Dlog4j2.formatMsgNoLookups=true` environment variable that turns off variable lookups, which was turning the network application to be not vulnerable to the Log4j exploit. It's also important to refer that this configuration uses a MongoDB database that supports the UniFi Network Application, where users are saved.

At first, an HTTPS connection with UniFi's dashboard could not be achieved as the default CA certificate was generated by an untrusted CA. Therefore, it was time to create a Certificate Authority responsible for issuing the signed public-key digital certificates associated with a predefined domain name. Turning the created CA into trusted in the target device made it possible to achieve an HTTPS connection.

The steps to load the certificates into the Docker container were as follows:

- 1) Map the certificates folder path to the `/usr/lib/unifi/cert` volume exposed by the container.
- 2) Insert in the certificates folder the PEM format SSL private key file corresponding to the server's SSL certificate under the name of `privkey.pem`.
- 3) Insert in the certificates folder the PEM format SSL certificate with the full certificate chain under the name of `fullchain.pem`.

The private key file is associated with UniFi's dashboard domain. The full chain file is simply a concatenation of the public-key certificates of the CA and the `example-domain.ui.com` server's domain.

Furthermore, Docker entry point scripts were changed to always reload SSL certificates inside UniFi's Docker container upon new executions. This was not the default behavior, as the Docker image was previously configured to issue a file with the hashes of the certificates and check for its existence. In such cases, SSL certificates were not reloaded.

After the above-mentioned SSL certificates are issued, the newly created root CA's public-key certificate is needed in the attacker machine to be considered trustworthy. This way, when the trainee visits the UniFi dashboard, the website appears with a legitimate HTTPS connection.

Every template file is within the scenario's setup folder. At first, the root CA's public-key certificate was copied into a special `ca-certificates` folder and run the `update-ca-certificates` command, which turns the

newly placed root CA certificate into system-wide trusted. So, every digital certificate signed by this new root CA will be deemed safe. After this, the Firefox browser needed to consider this CA safe. So, the `policies.json` file was copied into a special folder.

```
{
  "policies": {
    "Certificates": {
      "ImportEnterpriseRoots": true,
      "Install": [
        "ca.crt",
        "/setup/ca.crt"
      ]
    }
  }
}
```

Listing 7: Firefox's Policies File.

Listing 7 presents the policies file, which is read on every Firefox's new execution.

After loading the SSL certificates, the desired effect was obtained, an HTTPS connection when loading UniFi's dashboard. Still, there is a slight problem. When hitting the dashboard's web page for the first time, the initial wizard setup was shown. This was overcome by creating a Selenium script for this effect using Firefox's *WebDriver* to instruct the browser's behavior remotely. The tasks performed by Selenium can be summarized in:

- Visiting UniFi's web dashboard page.
- Setting administrator credentials for accessing UniFi's web page. These were specified in the YAML format as custom variables of the vulnerable service in the scenario's specific variables.
- Clicking several wizard setup buttons to move onto more advanced setup stages.

With these configurations set, it is time to move into the exploitation phase.

2) *Exploit*: The goal of the Log4j exploit on UniFi's software is to obtain a reverse shell, get the secret flag, and even leverage access to get the administrative credentials on the UniFi MongoDB instance as part of the post-exploitation process.

Initially, a little reconnaissance can be done by checking which ports are open by default in the victim machine using *nmap*. Indeed, port 8443 is open, which is where UniFi's interface is.

Then, the `https://example-domain.ui.com:8443` domain is accessed, and the login page from Fig. 6 is obtained.

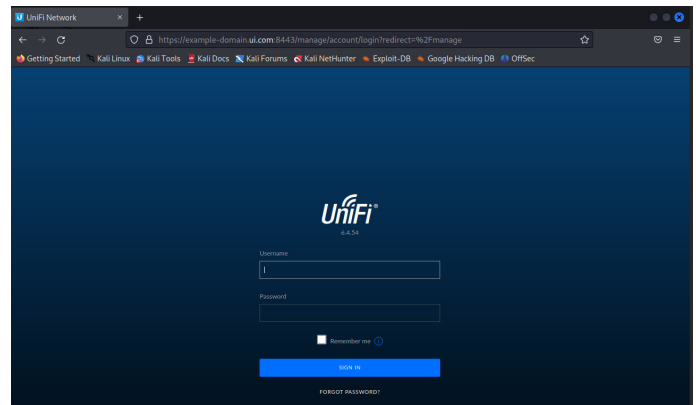


Fig. 6: UniFi's Initial Dashboard.

As mentioned earlier, the field to target needs to be logged by Apache Log4j using a malicious Java Naming and Directory Interface (JNDI) query. In this case, it is the `remember` field of the POST request.

Then, it is time to test if the web application is vulnerable to the Log4j attack. Firstly, the listening for connections on port 9999 using *netcat* with `nc -lnvp 9999` was done. From the POST request issued when submitting the login form, there is only the need to change the `remember` field to `${jndi:ldap://172.152.0.2:9999/whatever}` and issue the modified POST request. Notice `172.152.0.2` is the attacker machine's IP address.

As a connection is established on port 9999 from the vulnerable Log4j container, the web application is indeed susceptible to the exploit. The next step is to use the Rogue-JNDI GitHub tool to obtain a reverse shell on the target. Essentially, this tool sets up a malicious Lightweight Directory Access Protocol (LDAP) and HTTP server for JNDI injection attacks. When the tool first receives a connection from the vulnerable client to connect to the local LDAP server, it responds with a malicious entry containing a payload that will be useful to achieve a Remote Code Execution. The steps to build upon this foundation are to stage an LDAP Referral Server that will redirect the initial client request of the victim to an HTTP server where a secondary payload is hosted that will eventually run code on the target.

The procedures to set up the exploit include first using *netcat* to listen for inbound connections on port 4444 with the command `nc -lnvp 4444` and then following the steps:

- 1) Clone the Rogue JNDI tool and build the project into a JAR file using Maven.
- 2) Generate the Base64 payload that will run in the victim's server. It connects to the attacker machine on port 4444, redirecting both the standard input and standard output to the remote machine so the attacker can have complete control over the victim.
- 3) Running the Rogue JNDI tool to create malicious LDAP and HTTP servers with the command that will trigger a reverse shell.
- 4) Issue a cURL command with the malicious JNDI query and run the exploit.

As a result, a reverse shell is obtained in the initial *netcat* listener, meaning access was granted to the target server under the *unifi* user. Running a simple `ls -la` command, it is possible to see there is a weird file with the name “...” (3 dots). By opening it, the challenge flag `flag{l3ts_un1f1_every0ne_l0g4j}` is shown.

B. Ransomware Scenario

The Ransomware scenario is the first Windows-based scenario, opening the door to this new dissertation scope. The initial idea was to combine both Linux scenarios and Windows scenarios. There was a willingness to continue using containers to maintain consistency in the overall project. Still, since the development was based on a Linux host machine, and the underlying operating system resources and drivers used were also Linux-based, there was no way to create Windows containers. This happens because Docker is an OS-Level Virtualization, and the Docker daemon provides each container with the necessary kernel-level properties for it to be able to run. Due to this, Linux applications run on a Linux machine, and Windows applications run on a Windows platform. Still, there are exceptions in Windows due to the existence of *Linux Subsystem*, making it possible for a Linux container to run on Windows. With this in mind, the solution was to use Linux containers with KVM installed to run a Windows Vagrant box that would allow remote control. In the case of the Ransomware scenario, the malicious payload comes in the form of an executable (.exe) file, and having a Windows machine to run this script was the ideal situation.

This challenge distinguishes itself from the other scenarios because it is not attack-oriented. As such, there is no attacker machine and, therefore, no external network. Still, the final goal stays the same, which is to get the secret flag. This scenario is forensics-oriented in the sense that the trainee has to use a set of tools to debug the executable file, understand the consequences of executing the payload, and develop the reverse engineering skills necessary to get the flag.

1) Windows Vagrant Box Inside Linux Docker Container:

As mentioned, one cannot run Linux and Windows containers simultaneously using the same Docker daemon. The solution to overcome this problem was to install a Windows Virtual Machine inside a Linux container. From the Docker daemon’s perspective, all containers are Linux-based. Nonetheless, some of those containers run a hypervisor, on top of which there is a Windows Vagrant box. Ultimately, the goal is to configure and access the Windows machine through Remote Desktop (RDP). One may ask: *Why install a VM inside a container?* This may seem strange to many since installing the VM directly on the base OS is always possible without needing an extra container layer. However, running a VM inside a container has advantages in spinning up multiple identical Windows VMs, saving tremendous resources, mainly in terms of disk space.

When comparing a scenario where only a single VM runs directly on the base OS versus a scenario where the VM is

containerized, both situations consume similar resources. For instance, a VM that takes 30GB of disk space will take 35GB on a containerized setup. If six copies of a VM are run, the occupied disk space increases to 180GB, as each copy takes the exact amount of disk space. The situation slightly differs in the case of six copies of containerized VMs. In Docker, there are two distinct concepts: images and containers. Images turn out to be read-only and are the core of containers that are created from a read-only layer, the image. On top of this read-only layer, they add their own read-write layer, which differs between containers. Considering the example above, where the Docker image size is 35GB when creating six containerized VMs, each container will only vary in its read-write layer interacting with the read-only image. Assuming this read-write layer has a size of 10GB, all six containers have a combined size of 60GB on top of the 35GB Docker image, making a total of 95GB. To take this even further, one could consider using linked clones in Vagrant VMs in which new VMs only differing in disk images are created using the parent disk image belonging to a master VM.

RDP access was a desirable feature in these setups, but contrary to what happens in Linux, Windows containers cannot have a Desktop Environment. Instead, they are designed to run services and applications accessible using the PowerShell command line interface. Unlike Linux containers, where the Desktop Environment is an installable component, Microsoft ships Windows containers in a bundle directly with the OS. Microsoft published a set of known base images that form any Windows container’s base. For them, there is no installable Desktop Environment component, meaning even if using Windows containers, the issue of not having the possibility of remotely controlling the UI would be present.

The architecture of the Vagrant box can be seen in Fig. 7.

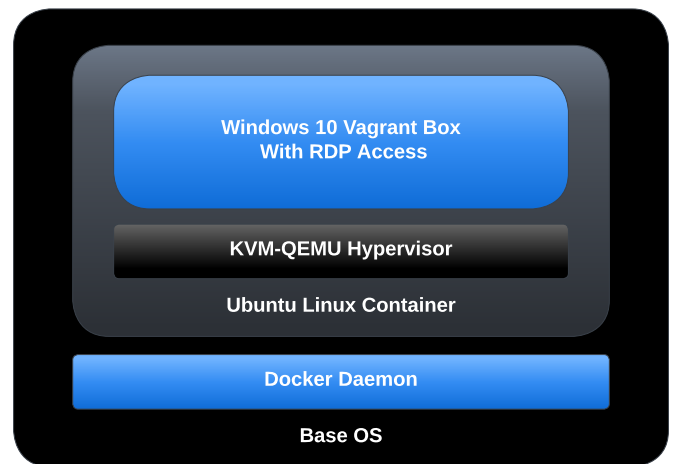


Fig. 7: Architecture Of Windows Vagrant Box Inside Docker Container.

This setup enabled a fully running Windows OS accessible through RDP and containerized and managed by Docker daemon. Five different technologies are worth mentioning:

- **Base Operating System**, that will be the main hosting platform.
- **Docker Daemon**, which will handle the final Docker image (*Ubuntu 18.04 Linux*) out of which a container will be spawned. The Docker image's main function is to run a hypervisor on which the Windows VM will run.
- **Hypervisor on the Docker Image (KVM-QEMU)**, which enables the installation and management of the Windows VM.
- **Windows VM**, the machine, a pre-packaged Windows 10 Enterprise Evaluation Vagrant box, that is available through RDP.

The first step is to build the Docker image with the hypervisor installed. For this, virtualization (VT-x) must be enabled in the BIOS settings to launch the VM. Then, in a *Ubuntu 18.04 Linux* image, firstly, the *QEMU-KVM* hypervisor package and *Libvirt*, an API library that manages KVM, are installed. Afterward, the `/dev/kvm` and `/dev/net/tun` devices in the host OS are mapped to the container, as well as the `/sys/fs/cgroup` directory, ensuring read-write permissions on it. Also, the container should run in privileged mode, meaning it can access almost all resources the host OS can. Another vital topic worth mentioning is the installation of Vagrant, which is necessary to run the Windows VM. At last, the respective *Vagrantfile*, which contains instructions on how to build the Vagrant box, is downloaded. The final Vagrant box size should be about 8.3GB.

Setting up the right *iptables* rules was a challenge. This is extremely important to ensure access to the RDP port on the Vagrant box from out of the container. By default, the Vagrant box configures firewall rules to allow access only from within the hypervisor container, meaning machines external to the hypervisor container do not have access to the Windows Vagrant box. As such, rules that redirect traffic from the base OS to the Vagrant box on RDP are needed. The logic followed is depicted in Fig. 8.



Fig. 8: Schema of Vagrant *iptables* Rules.

The inserted *iptables* rules on the hypervisor container concerning NAT and port forwarding from the host OS to the container were:

- Forward new TCP connections on ports 3389 (RDP), 5985 (PSRP HTTP), and 5986 (PSRP HTTPS) destined to the Windows VM.
- Add a “prerouting” rule that changes the destination packet address to the Windows VM on connections reaching ports 3389, 5985, and 5986.
- Add a “postrouting” rule that changes the source packet address to the hypervisor container on connections reaching ports 3389, 5985, and 5986.
- Forward established and related connections from and to the Windows VM.
- Reject every other traffic from and to the Windows machine. Notice the previous rules take precedence over this rule.

These rules are sufficient for establishing RDP and PSRP (PowerShell Remoting Protocol) connections. The former is a protocol for remote desktop access, while the latter is a protocol that runs over Windows Remote Management (WinRM). This remote management protocol uses a Simple Object Access Protocol (SOAP) based API for communication between the client and the server. Essentially, PSRP establishes remote sessions with the Windows machine, runs PowerShell commands and scripts on it, and receives the results back.

The PSRP traffic redirection rules denote how to forward traffic from Ansible instructions destined for Windows machines. After the creation of the Windows VM, there is the need to configure it, so the logic followed was the same as before, meaning using Ansible to configure the Vagrant box remotely. This way, commands issued from the base OS go through the Linux hypervisor container using the above-mentioned *iptables* rules and are redirected using NAT to reach the final target, the Windows VM box. This is possible using an SSH connection from the Ansible host machine to the hypervisor container, which will then redirect the traffic. Still, there are incompatibilities with these different remote access protocols between Linux and Windows: SSH and PSRP or WinRM.

A PSRP Ansible connector that connects to Windows-based machines using the PSRP protocol. A WinRM Ansible connector could as well be an option here. Still, PSRP offers the possibility to use a SOCKS5 proxy, which is suited for handling connections of Windows hosts sitting behind a bastion, in this case, the hypervisor machine. So, the current setup uses two different Ansible connectors: the Docker one that connects to the Linux containers and the PSRP one that connects to the Windows VM.

In the above paragraph, it was mentioned that the SOCKS5 proxy, which routes traffic back and forth between two distinct actors, acting as a middleman between the two. Packets going through this proxy are not modified nor encrypted, only in cases where traffic is encrypted through an SSH connection, as it currently happens, from the Ansible host to the bastion host. This SOCKS5 proxy is needed to forward WinRM commands to the bastion host. As mentioned, SSH creates incompatibility issues as it is only suited for remote access commands on Unix-like systems.

Fig. 9 shows a basic outline of the current configuration.

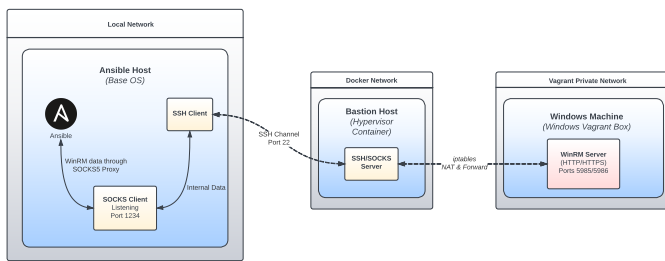


Fig. 9: Ansible Host, Hypervisor Container and Vagrant Box Architecture [16].

The network boundaries included in this setup include the following:

- **Ansible Host to the SOCKS Listener:**
 - The Ansible host forwards data using the WinRM payload encapsulated in a SOCKS packet.
 - A SOCKS5 proxy is set up in the Ansible host.
- **SOCKS Listener to the SSH client:**
 - Data from the SOCKS5 proxy is sent using an internal SSH channel.
- **SSH Channel:**
 - All data is encrypted using the SSH protocol.
- **SSH Server to the WinRM Listener:**
 - The bastion host, the hypervisor container, acts as the Ansible controller and sends the WinRM traffic to the Windows VM using port 5985 or 5986.
 - The WinRM service in the Windows VM sees the bastion host as the source of the communication and has no idea of the SSH and SOCKS implementation behind it.

Configuring the SSH proxy that exposes the SOCKS5 proxy to channel the WinRM requests through the bastion host is rather simple. The way to go is using SSH multiplexing with *ControlMaster*:

```
ssh -o "ControlMaster=auto" -o "
ControlPersist=no" -o "ControlPath=~/.
ssh/cp/ssh-%r@%h:%p" -CfNq -D
127.0.0.1:1234 kvm
```

Listing 8: SSH Proxy Exposing SOCKS5 Proxy.

The command in Listing 8 enables SSH multiplexing, which allows reusing an existing SSH connection to establish multiple sessions without needing to re-authenticate every time, saving resources. Without this, whenever a command is executed, the SSH client would need to establish a new TCP connection and a new SSH session with the remote host. A SOCKS5 proxy is also configured on port 1234. It creates a channel with the *kvm* host, an alias in the SSH configuration file for the hypervisor container, meaning the bastion host. After this is set up, the variables associated with the hypervisor

container in the Ansible environment must be configured, as shown in Listing 9.

```
"ansible_user": "administrator",
"ansible_password": "vagrant",
"ansible_connection": "psrp",
"ansible_psrp_protocol": "http",
"ansible_psrp_proxy": "socks5h://localhost
:1234"
```

Listing 9: Ansible Variables - Hypervisor Container.

In Ansible terms, the host machine issues commands to the Windows VM as if there is no bastion host in the middle. Here, the `ansible_psrp_proxy` variable pointing to the SOCKS5 proxy server is used, as specified in Listing 8. Any commands sent through it will be redirected to the Windows VM. Regarding the `socks5h` scheme, it means the DNS resolution is made in the bastion host, meaning the hypervisor container. Other variables such as `ansible_user` and `ansible_password` refer to the Windows VM's credentials. Regarding the `ansible_psrp_protocol`, port 5985 is used for HTTP connections.

2) *Scenario Construction*: The Ransomware scenario is based on the FireEye Flare-On Challenge of the 2016 edition and in the materials of the Malware Analysis and Incident Forensics course of the Sapienza Università di Roma. A Ransomware attack employs encryption to hold a victim's information at ransom. The target user or organization's critical data, which includes files, databases, or entire applications, are encrypted, and a ransom is demanded to provide access.

The Ansible construction of the scenario includes all the configurations presented in Section VI-B1 and some little extras:

- Copy of scenario files.
- Tool Installation:
 - **IDA Free Version** - Popular tool that allows users to debug, disassemble and decompile binary files.
 - **x64dbg** - Debugger and disassembler similar to IDA but designed mostly for Windows executables.
 - **Process Explorer** - Provides detailed information on processes, modules, handles, and threads running in Windows.
 - **Process Monitor** - Used for monitoring and capturing real-time system activity on Windows, including file system, registry, process, and network-related events.
 - **PeStudio** - Software analysis tool designed for examining files in the Windows PE (Portable Executable) format.
 - **Resource Hacker** - Tool that analyzes, modifies, and extracts resources in Windows executable files.

All these tools are installed by default in the Windows VM and are accessible to the trainee. As mentioned earlier, the VM is accessible through Remote Desktop, and the credentials for accessing it are *vagrant:vagrant* or *administrator:vagrant*.

3) *Reverse Engineering*: The process of obtaining a solution to the challenge requires going through the reverse engineering process. The following descriptions include figures of low-level Assembly code, which should also be the trainee's focus. At first, a basic static analysis will be performed. Afterward, code snippets will be explored.

Starting with some information *PeStudio* gives us. The binary is not packed, meaning the program is not obfuscated and compressed, making the analysis process more straightforward. This can be checked by the fact that the binary's sections have very low entropy. With a deeper inspection, a "Resource Section" is spotted, which shows an image using *Resource Hacker*. Using *PeStudio*, many API imports related to Microsoft's Crypto API can be observed, as well as other interesting imports associated with system parameters, loading resources, and locating files.

Moving on to the *IDA* analysis section, the challenge consists of two files: the malware executable and an encrypted file inside a folder named *briefcase*. The first block of code after the *main* function builds the "briefcase" Unicode string. The next system call is *SHGetFolderPathW* and is identified by the CSIDL parameter pointing to the desktop directory. Then, the binary checks if the length of the desktop directory path is smaller than 248. If it does, the execution moves forward, concatenating the "briefcase" string with the desktop path and storing it in a variable. This variable is then fed to the *CreateFileW* call, checking for the existence of a directory named "briefcase" in the Desktop. If not, execution terminates.

This malware sample contains a debugger trap because due to the fact that a dynamic analysis activity may end up closing a non-existent file handle and the debugging process immediately stops.

The next step is a *GetVolumeInformationA* call fetching volume C's serial number. This value is compared against 0x7DAB1D35h, meaning the malware targets a concrete machine that most likely does not match ours. If so, the execution terminates.

To move forward in the analysis, there is the need to patch the binary but keep the result of the subroutine with 0x7DAB1D35h, as this value will be later used in the execution. After the serial number check, the malware decodes a global variable using the above-mentioned serial number as a multi-byte XOR key. The final result string is "those-files-really-tied-the-fold-together." The next phase is related to starting the cryptography activities using Microsoft's Crypto API for file encryption. Firstly, the malware hashes the above-mentioned long string using SHA-1, deriving an AES-256 symmetric key. Later, it recursively enumerates every file in the "briefcase" directory and encrypts them. It uses Cipher Block Chaining (CBC) mode, being the Initialization Vector, the MD5 of the lower-cased name and the extension of each

file. After this value is set, two handles to the file are obtained: one for reading and one for writing. The read content goes through the *CryptEncrypt* function and is written back to the file in 16KB blocks.

If there is no file to be encrypted in the "briefcase" folder, the binary loads a resource, an image asking for a ransom, and sets it as the Desktop's Wallpaper.

Ultimately, the trainee must decrypt a previously encrypted file inside the briefcase folder to find the secret flag. Given that the malware uses AES symmetric encryption, one can take advantage of the fact that the key used for encrypting files is the same as the one used for decrypting them. So, one possibility that is not so straightforward for solving the challenge is to patch the binary by replacing the *CryptEncrypt* call for *CryptDecrypt* by modifying the sample's Import Address Table (IAT) statically or at runtime using a debugger. The other possible solution that is publicly available in the project's open-source repository is a Python script that computes the decryption key of the AES algorithm being the SHA-1 hash of the string "thosefilesreallytiedthefoldertogether," taking into consideration Microsoft's *CryptDeriveKey* inner workings, plus the MD5 hash of the lowercase of the filename and extension as the Initialization Vector. The decrypted file content is then unpadding. By applying such operations over the initially encrypted file inside the "briefcase" folder, it is possible to obtain the secret flag.

C. Active Directory Scenario

The last custom-made challenge is Windows-based and goes through the Microsoft Active Directory (AD) technology. AD complies with a database (or directory) and a set of services that connect users to the network resources they need.

1) *Scenario Construction*: The process of constructing this scenario is very similar to the structure presented in Section VI-B1. The main difference is the use of a Windows Server 2022 Evaluation box instead of a Windows 10 Enterprise Vagrant box. This change was added because the scenario focuses on building a vulnerable AD Domain Controller, which needs an underlying Windows Server. The other difference between the ransomware scenario and the AD one is the fact that now the attacker machine is sitting in the DMZ network, which simulates an attacker inside the organization's network. As with the ransomware scenario, there is no external network. The last added change concerned *iptables* rules from the bastion host to the Windows Server VM. Previously, only NAT and forwarding operations were performed with respect to the RDP and WinRM/PSRP connection, as the main goal was to access the remote VM using RDP and issue commands to it using the PSRP protocol. Now, as the challenge's focus is again attack-oriented, every type of traffic from the attacker machine to the bastion host should be redirected to the Windows Server VM. The only exception is SSH traffic because incoming connections to the bastion

host should not be redirected to the Windows Server VM, as they are destined to the hypervisor container and not to the Windows Server VM. With this configuration set, the trainee can issue attacks from the attacker machine to the Windows Server VM, knowing the traffic going through the bastion container will be correctly forwarded.

The development of a vulnerable AD Domain Controller was based on John Hammond's Active Directory Youtube series and currently existing GitHub sources.

The first configurations steps on the Windows Server VM are:

- 1) Change the Domain Controller's hostname to *DC01*.
- 2) Install Active Directory Services and create a domain named *xyz.com*.
- 3) Configure the DNS server as the Windows Server VM itself and create a reverse DNS zone.
- 4) Create a private network share controlled by Domain Administrators. Move the secret flag into it.
- 5) Allow Remote Desktop sessions to ordinary Domain users, as this is not enabled by default.
- 6) Change the administrator accounts default passwords.
- 7) Generate a vulnerable AD schema with a set of Domain users and the Domain groups they belong to, including information on the Domain Controller's local administrators. This schema is randomized on every new scenario execution.
- 8) Configure the Domain Controller with several kinds of vulnerabilities the trainee can explore using the previously generated vulnerable AD schema.

The vulnerable configuration steps include the following:

- Weaken the Domain Accounts password policy to allow weak passwords linked to user accounts.
- Create the AD Groups and Users and add them to the respective AD Group.
- Generate a vulnerable configuration to enable *Kerberoasting* attacks. Essentially, a service account with a weak password is created and it is specified that future ticket requests to this service account should use the easily crackable "RC4" Kerberos Encryption type.
- Configuration suitable for *AS-REP Roasting* attacks. A maximum of three user accounts is configured not to require Kerberos pre-authentication, enabling this kind of attack.
- Set a maximum of three AD user accounts as DNS Administrators.
- Set a maximum of three AD user accounts vulnerable to *DCSync* attacks.
- Disable SMB Signing which enables the existence of man-in-the-middle (MiTM) attacks on the SMB Server. The SMB protocol is typically used for sharing access to files, printers, and other resources across the network.

2) *Exploits*: This section intends to explore attacks on the Active Directory scenario. During the attack phase, we will use the following tools:

- **CrackMapExec**, a post-exploitation tool that helps automate assessing the security of large Active Directory domains. It supports several types of attacks, including various protocols such as SMB, LDAP, WinRM, and Kerberos.
- **Impacket**, a collection of Python modules for working with network protocols that are extremely useful for attacking Active Directory networks.
- **Bloodhound**, an Active Directory reconnaissance and attack management tool that depicts the AD network graphically and uses graph theory to identify hidden relationships, sessions, user permissions, and attack paths in a domain.
- **Mimikatz**, a tool that can exploit Microsoft's Authentication systems. It can perform attacks such as: *Pass the Hash*, *Pass the Ticket*, *Kerberoast Golden and Silver Tickets*, *DCSync* attacks, among others.

As in the Log4j scenario, the exploitation phase starts by doing some reconnaissance using, for instance, *nmap*. Indeed, port 22 of the hypervisor container, the machine responsible for redirecting traffic to the target Windows Server machine, is open.

The next step is to grab a set of commonly used Active Directory users and a subset of the *rockyou.txt* password dictionary to test if some AD user accounts are found. At first, the hypervisor container should be registered as the attacker machine's DNS server. Then, one can attempt to get some users using *CME* with: `crackmapexec ldap 172.100.0.40 -u users.txt -p '' -k`, where *172.100.0.40* is the bastion host, and the users file contains some of the commonly used AD usernames. The retrieved output tells us existing AD users, as well as information on the Domain Controller, for instance, the hostname (*DC01*) and the domain currently being targeted (*xyz.com*). With this information in mind, an entry in the */etc/hosts* file of the *dc01.xyz.com* domain can be mapped to the *172.100.0.40* IP address.

Moreover, a brute-force attack using both the users and passwords dictionary can be performed, again, using *CME* with `crackmapexec ldap dc01.xyz.com -u users.txt -p passwords.txt --continue-on-success | grep '[+]'`. With some luck, a match between the AD users and their respective passwords is obtained. Then, it is time to test the login in the Domain Controller using the credentials of a match with: `crackmapexec smb dc01.xyz.com -u USERNAME -p PASSWORD`. The command mentioned above plus an extra `--pass-pol` flag can be used to view the AD password policy. Other flags can be provided, such as the `--users` flag to view the currently existing users, a `--groups` flag to view the existing groups, or a `--computers` flag to view the devices that are part of the AD domain. All this information is helpful to perform similar brute-force attacks, given the gathered information on the used password policy, as well as the information on which AD users exist. Notice that with the credentials of an AD user, it is possible to have a Remote Desktop session linking

to the remote controller.

The next step is to use Bloodhound to view information on existing AD users and their groups. At first, one should configure Bloodhound and then use the *bloodhound-python* module, along with the previously fetched credentials of an AD user, and collect information on the Active Directory domain, using `bloodhound-python -u USERNAME -p PASSWORD -dc dc01.xyz.com -d xyz.com -c all`. This will generate a set of *JSON* files which should then be imported into Bloodhound.

Bloodhound provides a realistic view of several AD objects. Fig. 10 lists the Domain Users. Each of their attributes can be selected, providing crucial information about the groups they belong to, their unique identifiers, and other relevant information.

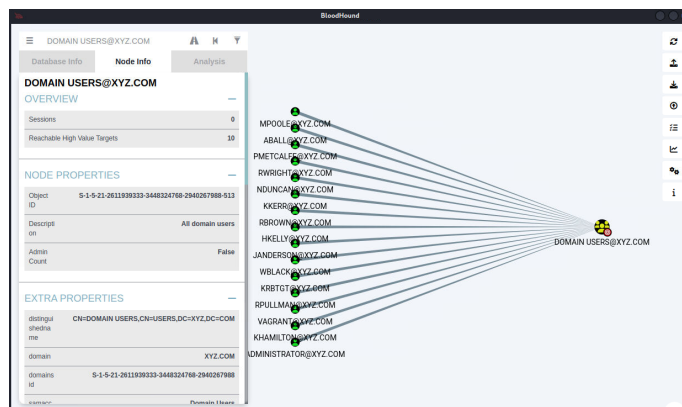


Fig. 10: Bloodhound Active Directory Users.

As a next step, one can gather information on the Domain Controller's administrators, which combine local and Domain Administrators, where a randomly selected account is given local administrator privileges for which the trainee can attempt to get the plaintext password using brute-forcing techniques.

Bloodhound also allows fetching AD user accounts vulnerable to *Kerberoasting*, *AS-REP Roasting*, and *DCSync* attacks. Finding dangerous permissions on Domain Users and Groups or legacy OS versions on computers belonging to the AD domain is also possible. Each may lead an attacker to a new attack path to obtain Domain Administrator privileges. At last, it allows the user to forge custom queries on the AD domain data, which turns this tool into one of the favorites for penetration testers.

Moreover, the trainee can use *Impacket* *WMIExec* module to perform lateral movement. This module allows executing commands on a remote system and establishing a semi-interactive shell on the remote host. As the current setup uses a bastion host, the connection is not directed to the Domain Controller, meaning there is the need to perform a tiny tweak on some module files. After all this, the trainee can run `impacket-wmiexec xyz.com/USERNAME:PASSWORD@dc01.xyz.com` to obtain a remote shell if using the local administrator account. It can then check the privileges of the current user

session with `whoami /priv` or try the same command with `impacket-psexec`. In this case, no remote shell is obtained because Windows flags this as a malicious operation. This attack works if the Domain Controller's Windows Defender is turned off. Simultaneously, the trainee can also check for RDP Desktop access using `impacket-rdp_check xyz.com/USERNAME:PASSWORD@dc01.xyz.com`.

Impacket also provides an SMB Client module to list the available network shares. With this tool, it is possible to observe the *SYSVOL* Domain Controller share listed, and also an *INTERNAL* share, on which only Domain Administrators have read access. To test this, the trainee can run `impacket-smbclient xyz.com/USERNAME:PASSWORD@dc01.xyz.com`. This share is where the secret flag is located.

Going more profound in the *AS-REP Roasting* theme, the *Kerbrute* tool that, similarly to *CrackMapExec* performs brute-force attacks on specific users, can be used. Furthermore, it identifies which accounts have pre-authentication disabled, meaning they are vulnerable to this attack. If a match between a user and password is found, it saves the Ticket Granting Ticket (TGT) ticket for each match found, which opens the door for *Pass the Ticket* attacks. To perform the *AS-REP Roasting* attack, the trainee can run `crackmapexec ldap dc01.xyz.com -u users.txt -p '' --asreproast out.txt`, where the `users.txt` file contains the AD user with pre-authentication disabled found using *Kerbrute*. This command captures the AS-REP response. Then, one can use *hashcat* to crack it and obtain the plaintext password with `hashcat -m18200 out.txt passwords.txt`.

It is time to enter the *Kerberoasting* world. The target here is service accounts, so each one must be enumerated. For this, the trainee should first brute-force the RID, the unique value representing an AD object. Then it can issue the command `crackmapexec smb dc01.xyz.com -u USERNAME -p PASSWORD -d xyz.com --rid-brute`. Then, it should create a file with the service account names and use the *GetUserSPNs Impacket script* with the credentials of an AD user account to get the Ticket Granting Service (TGS) tickets of the vulnerable service account. The command is as follows `python getUserSPNs.py xyz.com/USERNAME:PASSWORD -usersfile users.txt -output file hashes.kerberoast`, and finally, it is possible to crack the tickets using *hashcat* with `hashcat -m13100 --force -a 0 hashes.kerberoast passwords.txt`.

The next attack is *Pass the Ticket*, where the only premise is the fact that access is granted to the local administrator account in the Domain Controller and use *Impacket* to dump the NTLM hashes of the administrator account using `impacket-secretsdump -just-dc-ntlm xyz.com/USERNAME:PASSWORD@dc01.xyz.com`. The next technique is called *Overpass the Hash* because it uses an NT hash to obtain a Kerberos ticket that will be later used to impersonate a user. Then, to grab the TGT

ticket using the *Impacket GetTGT* module, the command `impacket-getTGT xyz.com/Administrator -hashes LMHASH:NTHASH` is used. Lastly, there is the need to export the *KRB5CCNAME* environment variable with the path of the TGT ticket and use *Impacket's* *WMIExec* module to get an Administrator shell on the Domain Controller using the *Pass the Ticket* attack. The trainee can also use *Impacket's* *SMBCClient* module to access the *internal* network share and grab the secret flag. This is the intended solution for solving the challenge.

The unintended solution uses *PsExec* to perform Local Privilege Escalation, which allows a non-admin process to escalate to *SYSTEM* if running *PsExec* with `.\PsExec64.exe -accept eula \\dc01 -s cmd` using the local administrator account on a Remote Desktop session. Lastly, the trainee should change the directory into the *internal* network share folder and print out the flag.

To perform the *Golden Ticket* attack, the trainee first grabs the NT hashes using *Impacket's* *Secrets Dump* module, as before, and the domain SID with `crackmapexec ldap dc01.xyz.com -u USERNAME -p PASSWORD --get-sid`. Then it can get a remote session using the local administrator account in the Domain Controller and, using *Mimikatz* in the remote machine, dump the *KRBtgt* account hashes by performing a *DCSync* attack with `lsadump::dcsync /domain:xyz.com /user:krbtgt`. Then, to craft the Golden Ticket, it needs the obtained Domain SID and the *KRBtgt* NT hash. With not just a Domain Controller but also workstations, the trainee should be able to run the attack and get administrator privileges in the Domain Controller.

Lastly, a DLL injection attack can be performed using a vulnerable AD account belonging to the *DNSAdmins* group. Essentially, this DLL is a reverse shell that connects to the attacker machine, and the threat actor should be able to obtain *SYSTEM* privileges on the Domain Controller machine.

The Domain Controller is vulnerable to a wide range of attacks. Some of them were presented above, but many others can be used to target the domain. As mentioned before, this scenario reveals itself as extremely useful as a way to provide the trainee with hands-on experiments for all the possible attacks.

VII. IMPORTED SCENARIOS

By now, the inner workings of the scenario construction process should be clear. The most complex scenarios were detailed in Section VI, but the project supports the addition of already existing scenarios. To keep the same logic followed so far was taken into consideration. As such, the final decision was to pick scenarios that were already based in Docker. Through profound research using platforms like CTFtime, several scenarios from the 2023 DiceCTF competition available on GitHub were chosen. A total of 23 challenges from the 2023 DiceCTF edition were imported. To enable this, a Python script was created to convert the scenarios from the format

they were published on GitHub to the one the framework understands.

Every imported DiceCTF scenario typically has a YAML file that is core to understanding the challenge. There, one can find information on the challenge's name, author, a short description, where to find the secret flag, the files that should be provided to the trainee so he can have a deeper understanding of the presented code, and the ports that the scenario's containers should expose to the exterior. With all this information in place, it was time to organize the folder structure for each scenario. The Python script first pulls the DiceCTF's GitHub repository and recursively looks for the challenges that use Docker. The ones that do not use it are, therefore, excluded. Other specific checks are also performed in this sense, for instance, challenges pinpointed to be excluded because their Docker setup is incompatible with the framework's setup, simply because they use Docker images that do not support the installation of *python3* and *iproute2* packages. Both these packages are required. The former is to be able to run Ansible playbooks and the latter to configure static routes between containers. Then, the script goes through the challenge mentioned above's YAML file and starts creating the scenario's variables file. As before, each scenario's vulnerable container is exposed via a domain, and the format followed by DiceCTF is *CHALLENGE_NAME.mc.ax*. In the variable's file, the necessary configurations in the DNS server are created, reverse proxy, and the edge router's port forwarding. As each scenario is attack-oriented, one must include the attacker machine in the external network of the setup. This will be the machine in control of the trainee. The following tasks involve setting up the custom scenario's structure, which is very similar for each scenario. The only things that change are the created containers and the domain of the exposed service. If the scenario has two or more containers, the container's names are mapped via the DNS server, as usually happens in Docker. For instance, if a scenario deploys containers *app* and *mongo*, the *app* container may attempt to access the database using the *mongo* domain. As such, the *mongo* string must be mapped to the *mongo* container's IP address.

VIII. SCENARIO EXTENSIBILITY

This topic intends to address the scenario extensibility theme. This topic is essential because it ensures the project's continuity regarding future updates. Scenario extensibility is closely related to what was presented in Section V-E. For this framework to support new scenarios, developers must understand the project's inner workings. This encompasses variable declarations for each scenario, namely DNS configurations for new domains, port forwarding issues, entry point setup, and representation of custom machines. Given the knowledge obtained so far, adding new scenarios to the current configuration should be relatively straightforward.

A new folder inside the *scenarios* folder should be created for each scenario. Then, a file named *challenge_vars.yml* should store the custom variables of the scenario. The rest

of the folder structure is dedicated to scenario's files. Every configuration should be specified on the *challenge_vars.yml* file.

Firstly, the DNS configuration, where information on each domain should be specified. This includes selecting the machines to which internal and external DNS requests should be mapped. Afterward, it is important to determine the Docker images that will be linked to the custom Docker containers of the scenario. This includes the Docker images' name, the path to reach the *Dockerfile*, and, optionally, possible arguments to the Docker image construction process. While explaining how the framework works, the idea of having a domain associated with the vulnerable machine was always taken into consideration because it is easier for the trainee to reach the vulnerable machine by a domain than by an IP address. While expanding the project's scenarios, the same line of logic was followed. As such, adding a reverse proxy to redirect each request to the appropriate Docker container is mandatory.

With this in mind, Docker image configuration for the proxy should also be included. The next step is to configure the Docker containers of the custom machines by specifying their name, the Docker image they are associated with, the groups they belong to, the Docker networks and assigned IP address, possible Docker volumes, information on where to find the DNS server, environment variables, and optional variables. The mandatory Docker containers include, again, the reverse proxy and the DNS server. An important note is that the reverse proxy configuration should consist of a set of variables later used in the NGINX configuration, as explained in Section V-F9. This includes specifying the domains to be mapped by the reverse proxy and the machines and ports the requests should be redirected. Notice that this setup allows load balancing configurations as several machines can be specified to the same domain.

The next topic concerns port forwarding configurations on the edge router. The goal of this configuration is to redirect requests from the attacker's machine located in the external network to the vulnerable service. But since there is a reverse proxy in between, requests are first redirected to it, which then forwards the packets to the final vulnerable service. This configuration allows us to configure the edge router's inbound port and configure the target machine and port to which packets should be redirected. In situations where the target machine is the reverse proxy, its destination port should be 443 (HTTPS), as specified in Section V-F9. In scenarios like Log4j, the connection between the attacker machine and the vulnerable service is not handled by a reverse proxy, meaning it can map packets directly to the vulnerable service.

Lastly, it is crucial to pay attention to the entry point scripting section. For configurations in Docker containers, there is a need to include a folder with a set of Jinja2 templating files that support the inclusion of variables specified in the YAML files. There is no need to use Jinja2 templates for configurations in the local machine. With each of these configurations, an *entrypoint.sh.j2* or an *entrypoint.sh* file should be created, as it will be the script that Ansible's actions

will trigger. Concerning the setup of the attacker machine, including the setup of the previously mentioned root CA, is needed so that the digital certificates are considered trusted across scenarios.

Regarding Windows-based scenarios, the only change is creating a hypervisor container that will host the Windows VM. To configure and provision the Windows VM, the developer may create a new set of Ansible playbooks that handle the configuration according to his will.

Having these considerations in mind and following the logic of the already presented scenarios, it is simple to expand the project to new scenarios.

IX. USER INTERFACE PANEL

For users that like to handle the scenario launching with the touch of a button, a user interface was created. It presents a listing of every scenario, as shown in Fig. 11.

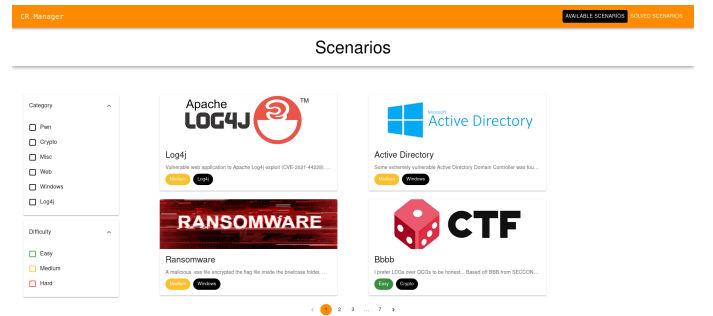


Fig. 11: User Interface Panel.

The UI panel consists of two primary tabs, one for all the available scenarios and one for the solved scenarios. Two sidebar filters are presented for the categories of the challenges and their difficulty. For each scenario, details associated with it, namely the name, author, category, difficulty, a short description, and the domains of the services to be targeted within the scenario, are visible to the user. Furthermore, there is a bottom panel where the user receives real-time feedback on the scenario's execution steps provided by Ansible via a *WebSocket*.

X. CLOUD DEPLOYMENT

One of the project's key aspects was the ability to run complex scenarios with the effort of a click or a command in the local machine. A UI panel was created to broaden the accessibility to users that want to keep it simple. However, the project focused on supporting remote deployments, namely, cloud deployments.

During the development, Microsoft Azure was the selected cloud provider. Two types of Azure VMs were created: a *Standard D2ads v5* and a *Standard E2bs v5*, featured with 2 vCPUs, 8GB and 16GB of RAM, respectively, and 128GB disk space due to the amount of memory the current cyber ranges take. The OS used was Ubuntu 22.04. Notice the selected VMs need to have KVM virtualization enabled;

otherwise, the setup of Windows-based scenarios will not work.

For every scenario deployment, the last steps include the installation of Tailscale in the attacker machine (*attackermachine*) and/or in the hypervisor container (*kvmcontainer*). After the installation of this tool, these machines are added to a Tailscale network as ephemeral nodes with a previously generated authentication key that one can use to sign in to the network. Even the Azure VM, where the project was later deployed, had Tailscale installed and joined the network using the same authentication key. With such a configuration set, accessing every machine belonging to the Tailscale mesh network was easy using *MagicDNS*. Through the host machine, which is also part of the Tailscale network, one can access the attacker machine using the domain `attackermachine.rhino-duck.ts.net`, the hypervisor container using `kvmcontainer.rhino-duck.ts.net`, and the remote machine using `local.rhino-duck.ts.net`. Notice `rhino-duck.ts.net` is Tailscale's network domain. With the remote deployment set in place, the trainee can use these domains to access every remote container without needing any port forwarding or firewall configurations set. This means it can access the attacker machine using VNC, Remote Desktop into the Windows VM, and even access the UI panel hosted in the remote machine. Notice that on every new scenario execution, either via the UI panel or the command line, the first step is attempting to remove possible attacker or hypervisor machines from the Tailscale network so that when new scenarios are deployed and other machines attempt to join the Tailscale network, there is no collision between domain names. Doing so ensures the correct machines are always targeted when accessing any of the domains mentioned above.

CONCLUSION

The current thesis touches on a multitude of topics and approaches. In this section, we go over each of these.

The literature review process revealed many cyber ranges built using old-case-driven approaches. To improve this process, a DevOps approach relying on Infrastructure as Code with Ansible to configure and provision cyber range scenarios based on Docker containers was used. Moreover, an enterprise-level network template was developed, considering various attack paths, enabling the trainee to solve a scenario in many ways. Randomization was addressed by changing the IP address of each container on every new scenario execution. Using Ansible, the created framework is capable of integrating a wide range of Docker-based scenarios supported by an enterprise-level network in a cost-effective manner. Creating custom scenarios with new environments related to world-known vulnerabilities, namely, Log4j, which haunted several companies across the end of the year 2021, meant a huge knowledge improvement. Furthermore, not only Linux operating system scenarios were visited. But with the aid of Vagrant, the project was expanded to Windows-based scenarios, which

introduced another degree of complexity. Among the developed scenarios, a Ransomware sample was used, replicating a threat that still haunts individuals and companies nowadays. A purposely vulnerable Active Directory Domain Controller was developed where trainees can test every sort of attack path. The journey of custom scenario creation does not finish here, as the framework was designed to allow the integration of new scenarios. To make the framework easy to manage, a UI panel was created. It works in sort of a CTF-like platform, where a user can launch scenarios, exploit them and submit the correct secret flag to mark them as solved. Finally, to make the framework flexible, the scenarios were deployed not only locally but to a remote machine and joined containers that needed to be accessible from the exterior to a Tailscale mesh network.

In conclusion, a straightforward open-source cyber range framework [3] that addresses the knowledge gaps the cybersecurity force needs to build upon was developed.

FUTURE WORK

The final outcome of the developed framework is quite interesting. However, many things can be considered for future work, namely:

Scenarios

Pursue further efforts in creating scenarios that use a wider range of network services, for instance, Mail servers, Intrusion Detection, and Prevention Systems. The possibility of using real hardware in some scenarios could also add value to the project. The ability to generate traffic from hypothetical users sitting in the network and logging the activities carried by an attacker to try to exploit a vulnerable service could also be meaningful from the trainee's point of view.

Scenarios Supporting Multiple Trainees

The current framework is suited for a single trainee only. A possible future improvement could be enabling scenario deployments to multiple remote machines so that numerous trainees can improve their skills in a style similar to a CTF competition.

Randomization

Several possible avenues of research were identified during the Literature Review (*cf.* Section IV), which were not fully explored. Regarding randomization, it was mentioned that every time a scenario is launched, the IP address of each container varies. The ability to experiment with several attack paths in certain scenarios was also taken into consideration. However, further developments can consider the randomization of user accounts, credentials, operating systems, other network configurations, and services. Even if considering the above-mentioned traffic generation, several types of traffic could be generated.

Docker Escape Concerns

This framework is tailored for running in controlled environments. Even in cloud deployments, access to externally accessible machines is guaranteed only if the machine trying to access them is part of the Tailscale network. The traffic exchanged between the mesh network is encrypted, and any device sitting outside it cannot access any machine or decrypt any traffic. Still, if by any means the trainee feels inspired to perform Docker escape attacks, it would be helpful to explore some countermeasures to prevent the trainee from accessing the host system.

ACKNOWLEDGMENT

This study was made possible in part with support from the Faculty of Engineering of the University of Porto, from Rolando Martins, partially funded by POCI-01-0247-FEDER-047264 (Theia) financed by Fundo Europeu de Desenvolvimento Regional (FEDER), through COMPETE 2020 and Portugal 2020 and from Carlos Novo, funded by FCT, Fundação para Ciência e Tecnologia, through the individual research grant 2021.08532.BD.

REFERENCES

- [1] Rebecca Acheampong et al. "Security Scenarios Automation and Deployment in Virtual Environment using Ansible". In: 2022 14th International Conference on Communications (COMM). 2022, pp. 1–7. DOI: 10.1109/COMM54429.2022.9817150. URL: <https://doi.org/10.1109/COMM54429.2022.9817150>.
- [2] Another Log4j on the fire: UniFi. Available at <https://www.sprocketsecurity.com/resources/another-log4j-on-the-fire-unifi>. Last accessed in May 2023.
- [3] Cyber Range Framework: GitHub. Available at <https://github.com/2dukes/Cyber-Range-Framework>. Last accessed in August 2023.
- [4] Razvan Beuran et al. "Integrated framework for hands-on cybersecurity training: CyTrONE". In: Computers & Security 78 (2018), pp. 43–59. ISSN: 0167-4048. DOI: <https://doi.org/10.1016/j.cose.2018.06.001>. URL: <https://www.sciencedirect.com/science/article/pii/S0167404818306527>.
- [5] Francesco Caturano, Gaetano Perrone, and Simon Pietro Romano. "Capturing flags in a dynamically deployed microservices-based heterogeneous environment". In: 2020 Principles, Systems and Applications of IP Telecommunications (IPTComm). 2020, pp. 1–7. DOI: 10.1109/IPTComm50535.2020.9261519. URL: <https://doi.org/10.1109/IPTComm50535.2020.9261519>.
- [6] Wenliang Du. "SEED: Hands-On Lab Exercises for Computer Security Education". In: IEEE Security & Privacy 9.5 (2011), pp. 70–73. DOI: 10.1109/MSP.2011.139. URL: <https://doi.org/10.1109/MSP.2011.139>.
- [7] Bernard Ferguson, Anne Tall, and Denise Olsen. "National Cyber Range Overview". In: 2014 IEEE Military Communications Conference. 2014, pp. 123–128. DOI: 10.1109/MILCOM.2014.27. URL: <https://doi.org/10.1109/MILCOM.2014.27>.
- [8] Massimo Ficco and Francesco Palmieri. "Leaf: An open-source cybersecurity training platform for realistic edge-IoT scenarios". In: Journal of Systems Architecture 97 (2019), pp. 107–129. ISSN: 1383-7621. DOI: <https://doi.org/10.1016/j.sysarc.2019.04.004>. URL: <https://www.sciencedirect.com/science/article/pii/S1383762118304442>.
- [9] Tommy Gustafsson and Jonas Almroth. "Cyber Range Automation Overview with a Case Study of CRATE". In: Secure IT Systems: 25th Nordic Conference, Nord-Sec 2020, Virtual Event, November 23–24, 2020, Proceedings. Berlin, Heidelberg: Springer-Verlag, 2020, pp. 192–209. ISBN: 978-3-030-70851-1. DOI: 10.1007/978-3-030-70852-8_12. URL: https://doi.org/10.1007/978-3-030-70852-8_12.
- [10] Hetong Jiang, Taejun Choi, and Ryan K. L. Ko. "Pandora: A Cyber Range Environment for the Safe Testing and Deployment of Autonomous Cyber Attack Tools". In: Security in Computing and Communications. Ed. by Sabu M. Thampi et al. Singapore: Springer Singapore, 2021, pp. 1–20. ISBN: 978-981-16-0422-5. DOI: 10.1007/978-981-16-0422-5_1. URL: https://doi.org/10.1007/978-981-16-0422-5_1.
- [11] Ryotaro Nakata and Akira Otsuka. "CyExec*: A High-Performance Container-Based Cyber Range With Scenario Randomization". In: IEEE Access 9 (2021), pp. 109095–109114. DOI: 10.1109/ACCESS.2021.3101245. URL: <https://doi.org/10.1109/ACCESS.2021.3101245>.
- [12] G. Perrone and S. P. Romano. "The Docker Security Playground: A hands-on approach to the study of network security". In: 2017 Principles, Systems and Applications of IP Telecommunications (IPTComm). 2017, pp. 1–8. DOI: 10.1109/IPTCOMM.2017.8169747. URL: <https://doi.org/10.1109/IPTCOMM.2017.8169747>.
- [13] Cuong Pham et al. "CyRIS: A Cyber Range Instantiation System for Facilitating Security Training". In: Proceedings of the 7th Symposium on Information and Communication Technology. SoICT '16. Ho Chi Minh City, Vietnam: Association for Computing Machinery, 2016, pp. 251–258. ISBN: 9781450348157. DOI: 10.1145/3011077.3011087. URL: <https://doi.org/10.1145/3011077.3011087>.
- [14] Z. Cliffe Schreuders et al. "Security Scenario Generator (SecGen): A Framework for Generating Randomly Vulnerable Rich-scenario VMs for Learning Computer Security and Hosting CTF Events". In: ASE @ USENIX Security Symposium. 2017.
- [15] Michael Thompson and Cynthia Irvine. "Labtainers Cyber Exercises: Building and Deploying Fully Provisioned Cyber Labs That Run on a Laptop". In: Proceedings of the 52nd ACM Technical Symposium on Computer Science Education. SIGCSE '21. Virtual Event, USA: Association for Computing Machinery, 2021, p. 1353. ISBN: 9781450380621. DOI: 10.1145/3408877.3432490. URL: <https://doi.org/10.1145/3408877.3432490>.
- [16] Windows Host Through SSH Bastion on Ansible. Available at <https://www.bloggingforlogging.com/2018/10/14/windows-host-through-ssh-bastion-on-ansible/>. Last accessed in May 2023.